
Active Object Localization

— Reinforcement Learning Final —
Project 2021

Meet Us

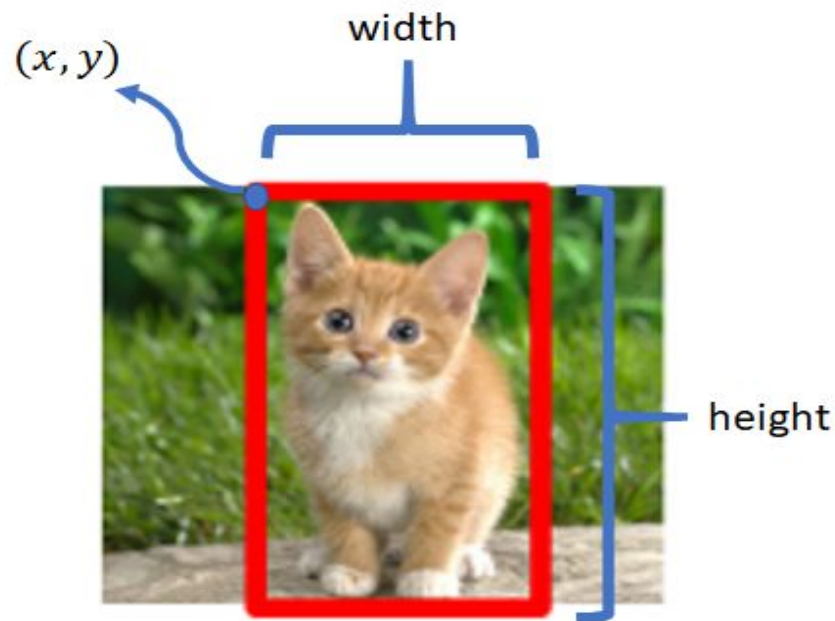
- Mohammed Deifallah: Mohammed.Deifallah@skoltech.ru
- Alexander Vikhorev: Alexander.Vikhorev@skoltech.ru

Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Motivation

- Collection of CV tasks.
- RCNN Family
- YOLO Family
- New Level of RL



Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Related Work

- Based on [Active Object Localization with Deep Reinforcement Learning](#)
- Per-class RL agents
- **MDP** (Markov Decision Process): **A** (actions), **S** (state), **R** (reward function)
- ϵ -greedy policy



Related Work (Method)

State

o (4096)

h (9 x 10)

Actions

Horizontal moves



Right



Left

Vertical moves



Up



Down

Scale changes



Bigger



Smaller

Aspect ratio changes



Fatter



Taller



Trigger

Reward

$$R_a(s, s') = \text{sign}(\text{IoU}(b', g) - \text{IoU}(b, g))$$

Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Dataset

- PASCAL VOC 2007-2012 (only 2007 😊)
- 20 classes (only 5 😊)
- PyTorch (only offline 😊)

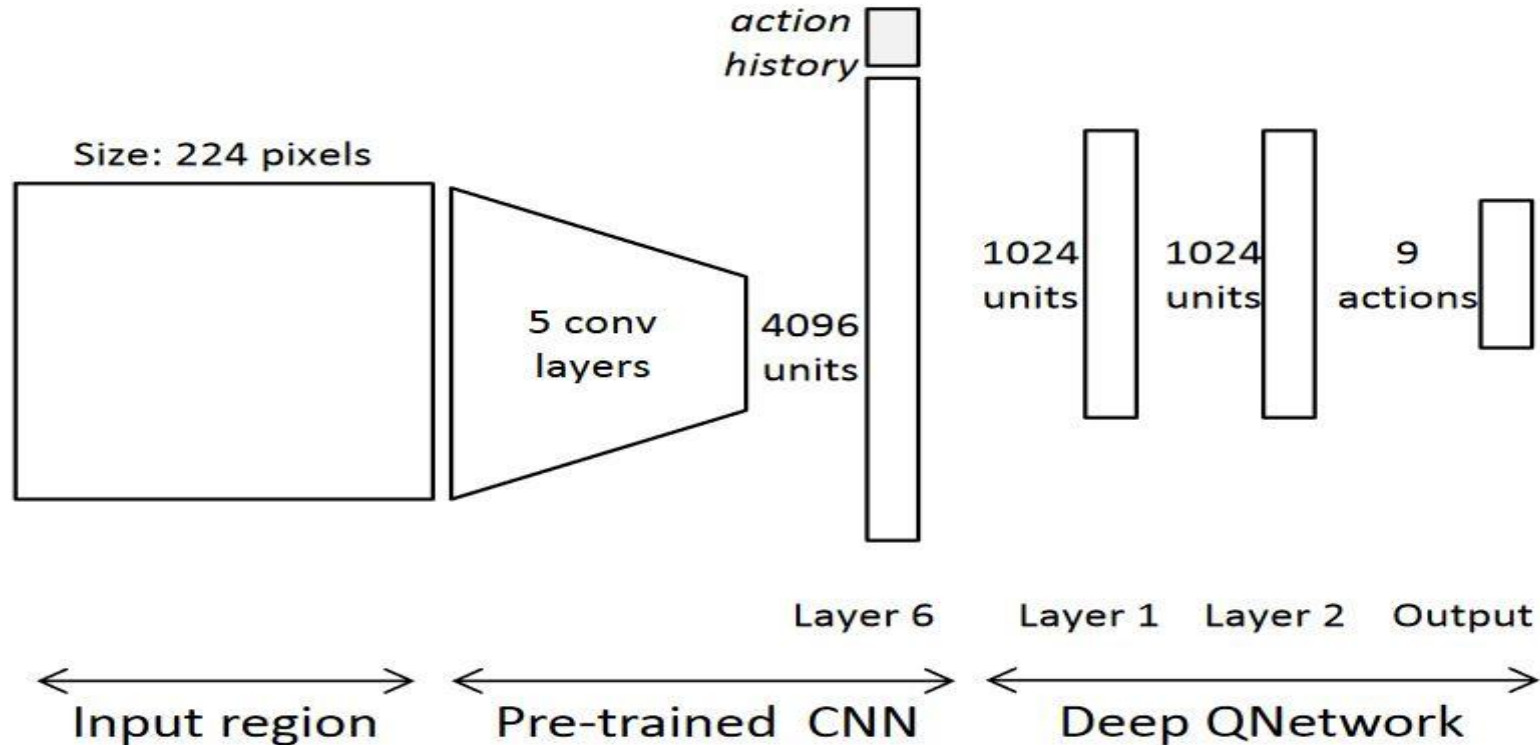


Source: <http://host.robots.ox.ac.uk/pascal/VOC/>

Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Model Architecture



Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Training

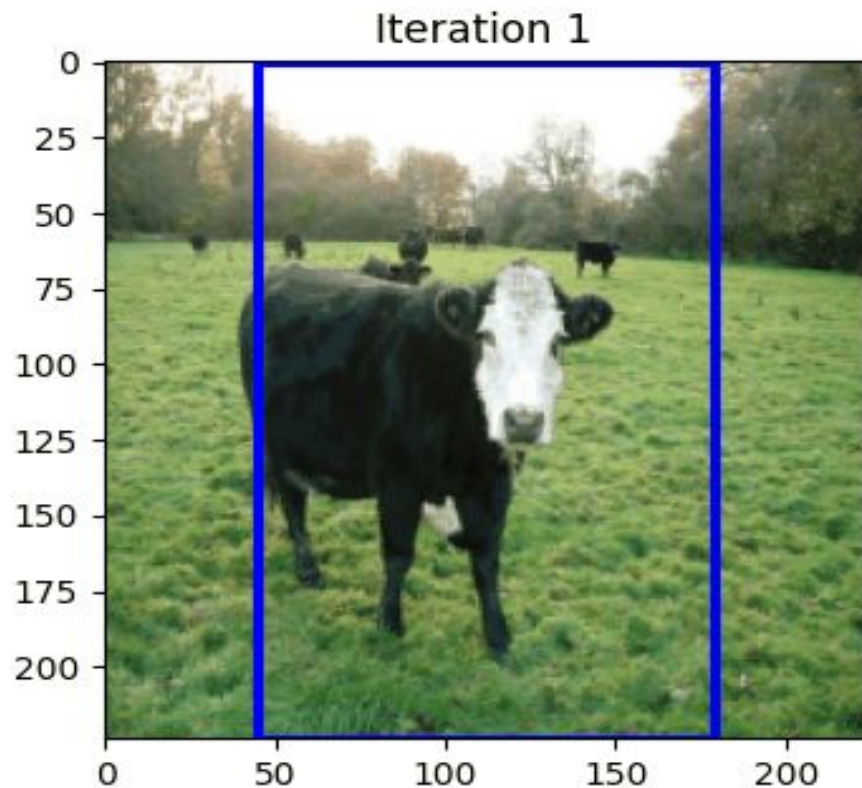
- Pascal VOC 2007 training set
- Input: 224 x 224 image
- $\alpha = 0.2$
- $\varepsilon = 1.0 \rightarrow 0.1$
- $\eta = \text{sign}(3.0)$
- Replay Memory
- 15 episodes

Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Evaluation

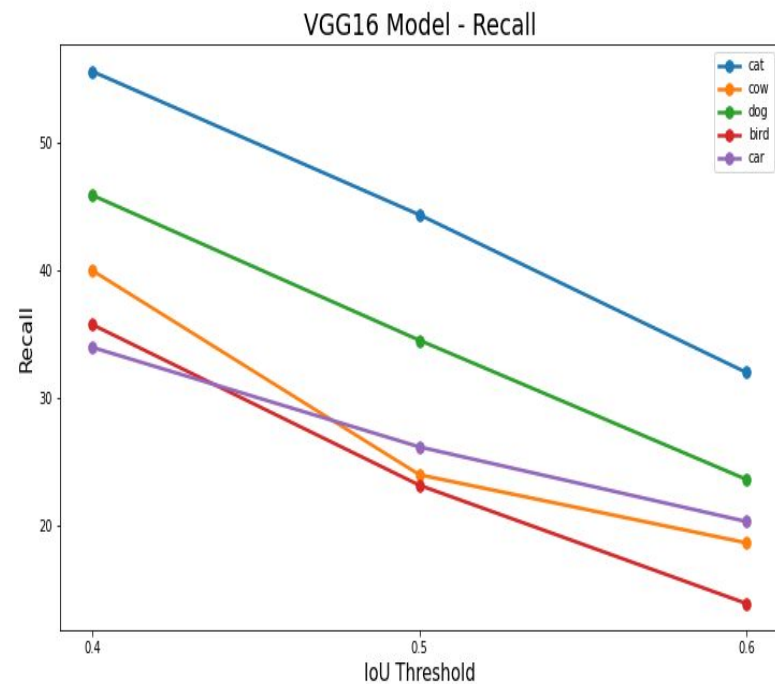
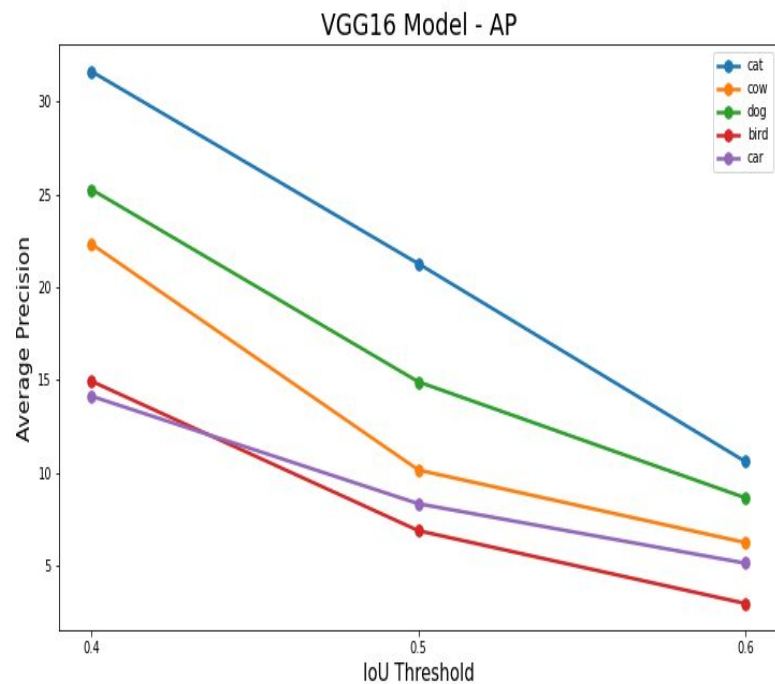
- PASCAL VOC 2007 validation set
- AP (Average Precision) & Recall
- Max. 40 iterations



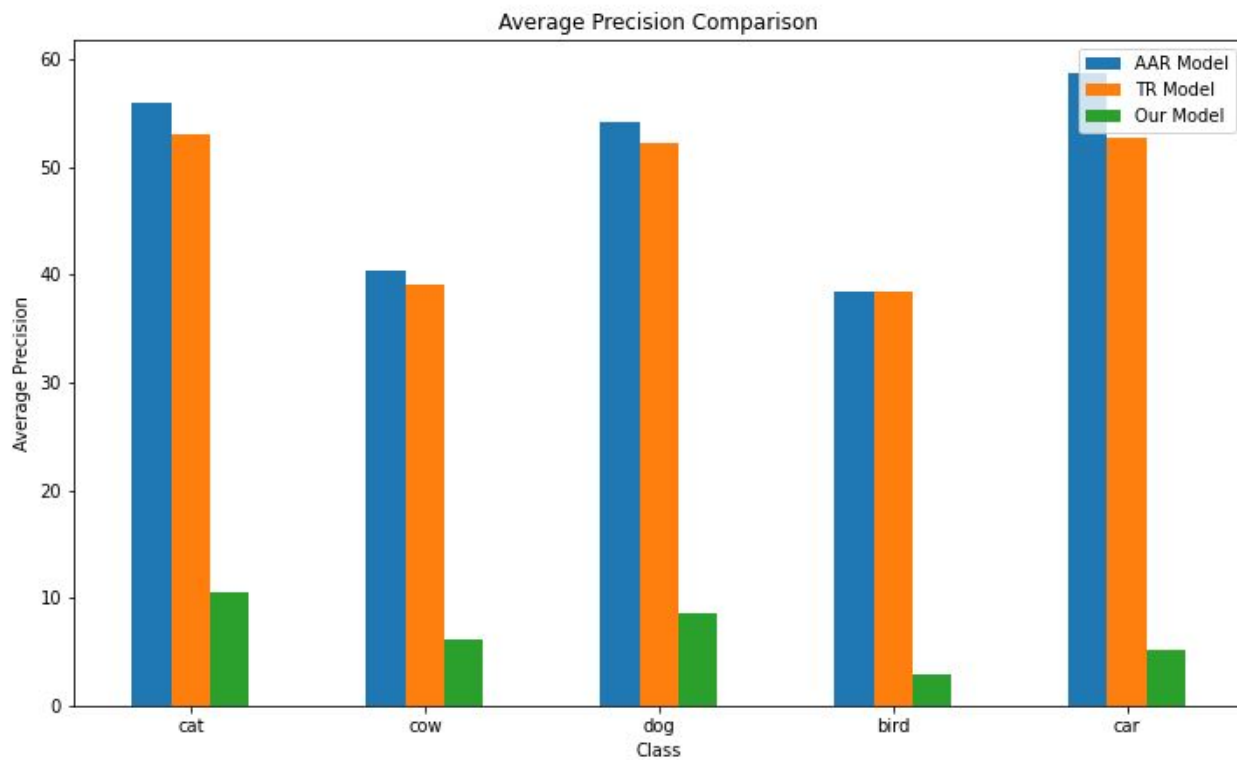
Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Results



Results



Results (GitHub Repo)

Code Overview

- `Training.ipynb` is used to reproduce the training process of the model.
- `Testing.ipynb` is used to reproduce the testing process of the model and visualize some examples of localization.
- `Plotting.ipynb` is used to plot all graphs and charts shown above using `media` folder.
- `media` is a folder to save all examples of localization and graphs.
- `models` is a directory which is needed to keep the saved models weights after training. This is an [example](#) of the so-called folder.
- `utils` is a directory containing the following files:
 - `agent.py` : a wrapper for the per-class agent that contains the whole components of RL (ϵ -greedy policy, reward, ... etc).
 - `models.py` : a wrapper for the two main modules of the network: *Feature Extractor* and *DQN*.
 - `dataset.py` : a separate file for reading the dataset (train and val).
 - `tools.py` : a collection of useful functions, such as computing the metrics or extracting a specified class from the dataset.

Source: <https://github.com/Mohammed-Deifallah/Active-Object-Localization>

Agenda

- Motivation
- Related Work
- Dataset
- Model Architecture
- Training
- Evaluation
- Results
- Conclusion

Conclusion

- Linear SVM should be added to the model.
- IoU threshold > 0.6 has negative effect.
- Different pre-trained networks (e.g. ResNet50)

Bonus slides : replacing DQN with PPO

- All environment related logic has been factored out and wrapped by gym.Env interface
- Allows to plug-in any available out-of-box algorithm implementation which works with Open AI Gym (Stable Baselines, Tianshou, Elegant RL, Lightning Bolts, ...)
- PPO implementation from Stable Baselines 3 has been chosen for its ability to handle both discrete and continuous inputs as well as learning performance
 - Actor-Critic scheme - learns both the policy and value function
 - On-policy algorithm - updates the model episode by episode based on the current exploration of the agent

Bonus slides : PPO (objective function)

$$L_t^{CLIP+VF+S}(\theta) = \hat{\mathbb{E}}_t \left[L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t) \right]$$

The diagram illustrates the components of the PPO objective function. It features four horizontal lines: a purple line under $L_t^{CLIP}(\theta)$, a green line under $-c_1 L_t^{VF}(\theta)$, and a red line under $+c_2 S[\pi_\theta](s_t)$. A blue line connects the purple and green lines, and another blue line connects the green and red lines. A red line connects the red line to the text 'Add an entropy bonus to ensure sufficient exploitation.'

c1 and c2 are coefficients.

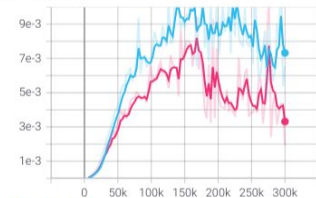
Squared-error value loss: $(V_\theta(s_t) - \underbrace{V_t^{\text{targ}}}_{\text{wavy line}})^2$

Add an entropy bonus to ensure sufficient exploitation.

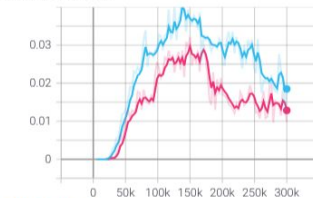
- 3 parts : PG loss , VF loss, Entropy loss
- Figure source : <https://spraphul.github.io/blog/RL5>

Bonus slides : PPO (training metrics)

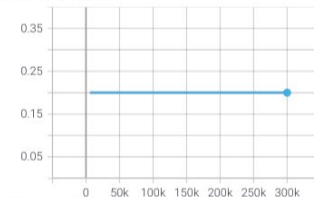
train/approx_kl
tag: train/approx_kl



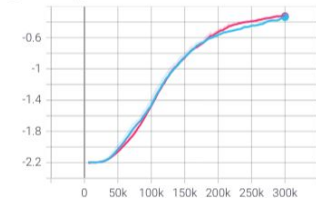
train/clip_fraction
tag: train/clip_fraction



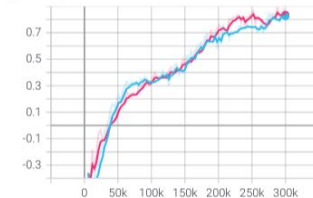
train/clip_range
tag: train/clip_range



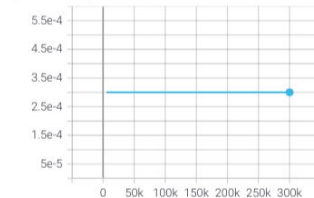
train/entropy_loss
tag: train/entropy_loss



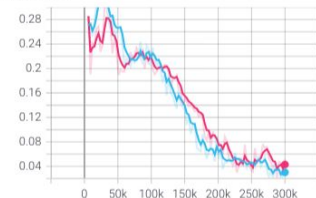
train/expained_variance
tag: train/expained_variance



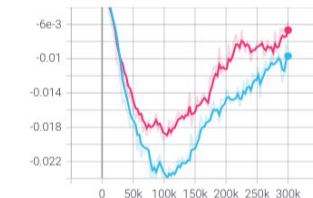
train/learning_rate
tag: train/learning_rate



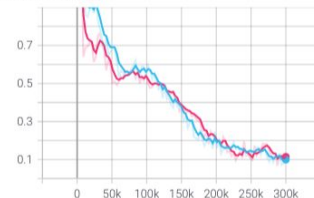
train/loss
tag: train/loss



train/policy_gradient_loss
tag: train/policy_gradient_loss



train/value_loss
tag: train/value_loss



KL-divergence - how different is the updated policy from the previous policy (we don't want this metric to be high since we want gradual change of the policy)

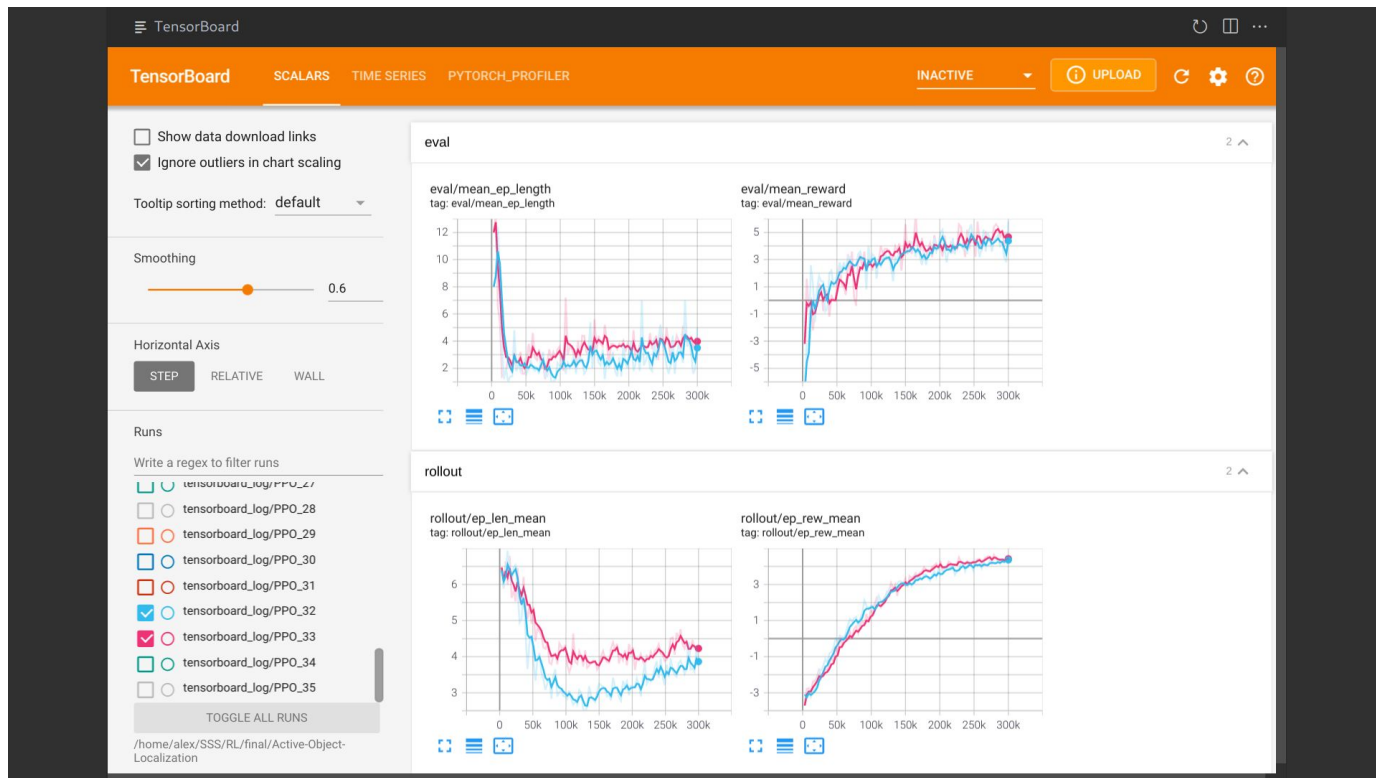
Clip fraction - Fraction of times the clip range hyperparam is used. PPO clips the new policy within the clip range of the old policy (stable learning).

Clip range - defined by hyper-param (usually 0.1 - 0.2)

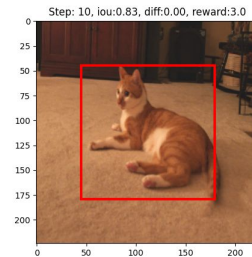
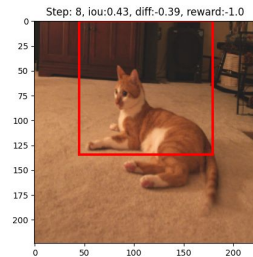
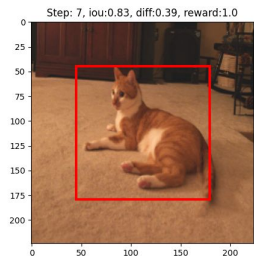
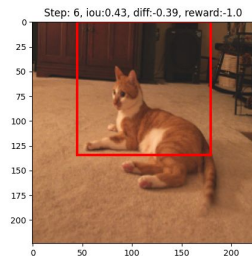
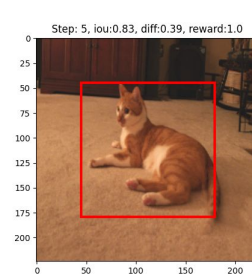
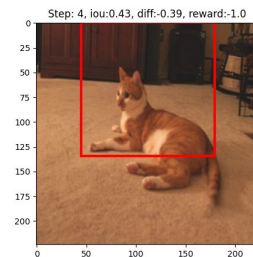
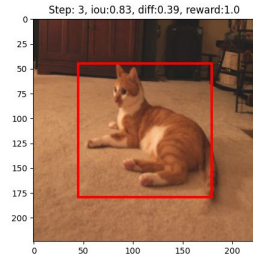
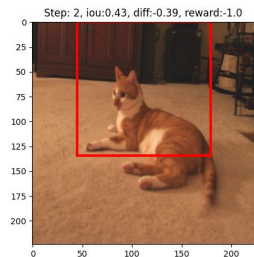
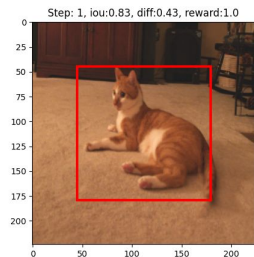
Entropy loss increases (in the objective function we reward high entropy and entropy decreases over time (in the beginning favoring exploration over exploitation and vice versa in the end))

General train loss, policy gradient loss and value function loss decreases over time

Bonus slides : PPO (mean episode reward)



Bonus slides : PPO (results)



Thanks!

Questions?